

# Closed-Loop AI-Native Threat-Driven SSDLC System

With Autonomy-Weighted Risk Scoring &  
Automated Threat-to-Test Transformation

*A Security Operating System for Autonomous AI Systems*

Closed-Loop

AI-Native

Threat-Driven

Autonomy-Scored

Self-Improving

# Traditional SSDLC Cannot Secure Autonomous AI Systems

Existing security models assume deterministic software behavior

## Non-Deterministic Behavior

Outputs vary per inference — traditional testing assumes repeatable results

## Tool Autonomy

Agents invoke external tools with real-world side effects beyond developer control

## Memory Mutation

Persistent memory enables cross-session contamination and state poisoning

## Multi-Turn Reasoning

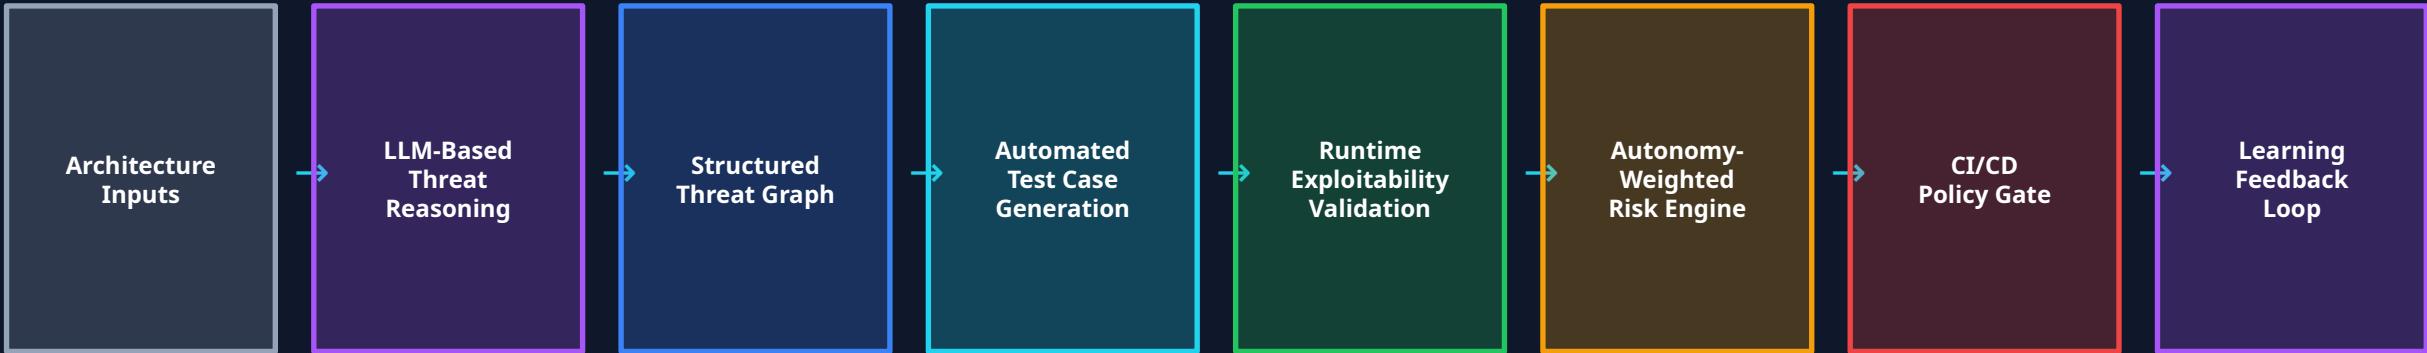
Complex reasoning chains create emergent attack surfaces invisible to static analysis

## THE CORE GAP

No existing system dynamically links threat modeling, exploit validation, autonomy scoring, and CI/CD enforcement in a closed loop.

# Closed-Loop Threat-Driven AI Security Engine

*Core Patent System Architecture*



**CONTINUOUS FEEDBACK LOOP:** Runtime signals regenerate threat models, rules, and test cases

## KEY INNOVATIONS

Threat outputs automatically generate executable dynamic test scenarios

Runtime exploitability signals modify static analysis rules

Adversarial testing results auto-update agent guardrails

AI Autonomy Index dynamically weights risk computation

The novelty is in the closed-loop integration logic — not the individual components

# Dynamic Risk Computation for Autonomous AI Systems

TRADITIONAL: Risk = Impact x Likelihood x Exposure

MISSING: AI Autonomy

X

**Risk = Impact x Likelihood x Exposure x Autonomy Index**

Autonomy Index =  $f($  )

Tool Write Access

Can the agent modify external systems?

Irreversible Action Capability

Can actions be rolled back?

Memory Persistence

Does the agent retain cross-session state?

Human-in-the-Loop Presence

Is human approval required?

Cross-Agent Interaction Scope

Can it influence other agents?

Autonomy increases blast radius even if exploit probability is unchanged

CLAIMABLE INNOVATION — No current risk model incorporates autonomous agent capability into security scoring

# Same Vulnerability, Different Risk

Vulnerability: Prompt Injection via concatenated user input | Same CVSS: 8.1

## Read-Only Chatbot

Autonomy Level: LOW

- No tool write access
- No persistent memory
- Human approves all outputs
- Single-agent, no cross-system reach

RISK: 4.1 — ALLOWED

## Database Write Agent

Autonomy Level: HIGH

- Full database write access
- Persistent memory across sessions
- No human-in-the-loop
- Multi-agent communication enabled

RISK: 9.2 — BLOCKED

**Same vulnerability, same CVSS — fundamentally different real-world risk**

The Autonomy Index captures what static severity scores cannot — the agent's capacity for irreversible, unsupervised action

# Automated Threat Vector to Executable Attack Translation

**Today: Threat modeling outputs documents — not executable security tests**

## LLM Threat Reasoning Engine Outputs

<b>Attack Path</b>	Full exploitation chain from entry to impact
<b>Target Asset</b>	Specific component, endpoint, or agent
<b>Required Access</b>	Authentication level and privileges needed
<b>Exploit Pattern</b>	Classified attack technique and payload type

## Auto-Generated Test Outputs

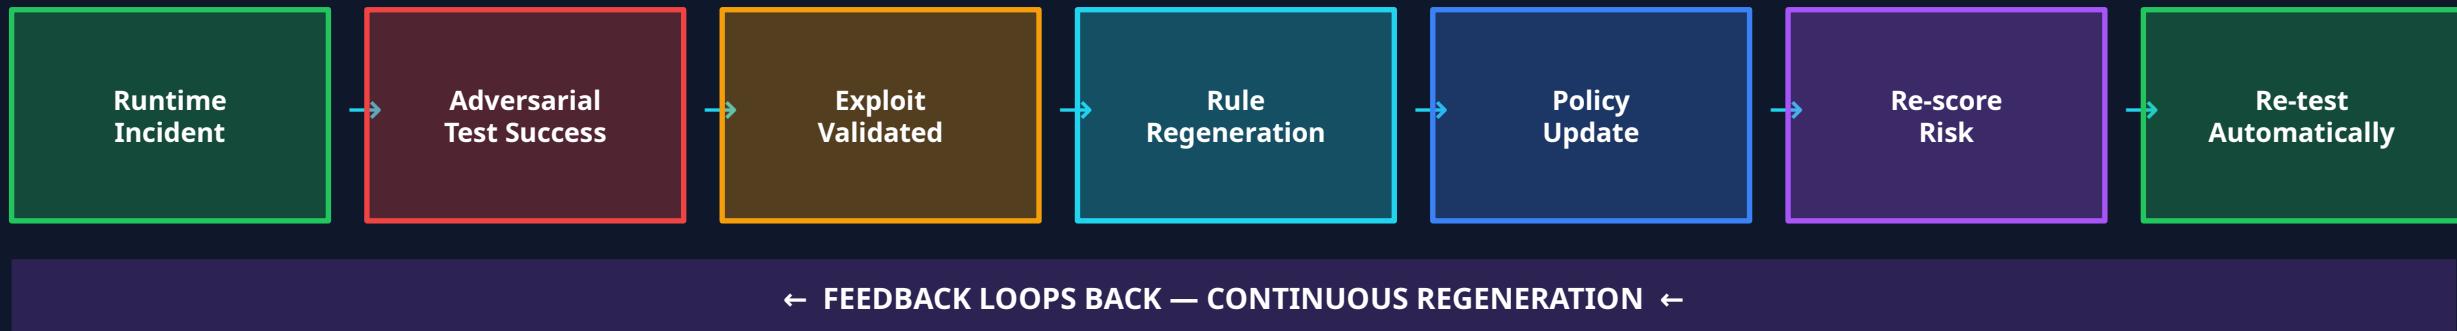
<b>Browser-Based Attack Script</b>	End-to-end UI attack simulation
<b>LLM Adversarial Test Case</b>	Jailbreak and injection validation
<b>API Fuzz Scenario</b>	Mutation-based API exploitation
<b>Agent Tool Abuse Test</b>	Permission escalation validation

**Threat Vector → Structured Graph → Test Template → Executable Attack → Result Feedback**

## CLAIMABLE INNOVATION

The system reasons about threats using LLM intelligence, then automatically generates executable attack simulations that validate real exploitability — a non-obvious automation beyond existing tooling.

# Closed-Loop Self-Improving Security System



**KEY CLAIM: Runtime exploitability updates static analysis rules automatically**

## What Gets Regenerated Automatically:

<b>Static Analysis Rules:</b>	Detection patterns evolve from confirmed runtime exploits	<b>Dynamic Test Scenarios:</b>	Attack simulations updated with newly discovered bypass methods
<b>Agent Guardrail Constraints:</b>	Prompt boundaries, tool restrictions, output filters auto-updated	<b>Autonomy Scoring Thresholds:</b>	Risk weights recalibrated from real-world incident impact data
<b>CI/CD Policy Gate Thresholds:</b>	Blocking and approval thresholds adjusted from enforcement outcomes		

**Full automation — no manual rule updates, no human bottleneck — the loop is the invention**

# Patent Claims & Differentiation

CLAIM 1

## A Closed-Loop AI-Native SSDLC System

A method and system comprising continuous threat modeling, automated security testing, runtime exploit validation, and feedback-driven rule regeneration operating as a unified closed-loop pipeline for securing autonomous AI systems.

CLAIM 2

## An Autonomy-Weighted Dynamic Risk Engine

A computational method for dynamically scoring security risk by incorporating an Autonomy Index — a composite of tool write access, irreversible action capability, memory persistence, human-in-the-loop presence, and cross-agent interaction scope — as a multiplicative risk factor.

CLAIM 3

## Threat-to-Test Automated Transformation

A method comprising: parsing architectural representations, generating structured threat vectors via LLM-based reasoning, translating vectors into executable dynamic test cases (browser attacks, adversarial prompts, API fuzz, agent abuse tests), and updating the threat graph from runtime exploitability signals.

### vs. Traditional DevSecOps

Static severity scoring →

**Autonomy-weighted dynamic scoring**

Manual threat modeling →

**LLM-generated structured threat graph**

Post-deployment validation →

**Continuous exploitability-driven regeneration**

**Filing Strategy: Narrow claims on integration logic — the novelty is the closed-loop architecture**